

---

# Locality and The Fast File System

# First File System

- “old UNIX file system” by Ken Thompson



super block

- simple
- supported files and the directory hierarchy



Kirk McKusick


- The problem: performance was terrible.

- Performance started off bad and got worse over time, to the point where the file system was delivering only 2% of overall disk bandwidth! [Kirk McKusick '84]

100MB/s → 2MB/s

- Treated the disk like it was a random-access memory; data was spread and thus had real and expensive positioning costs. → seek time
- For example, the data blocks of a file were often very far away from its inode

inode와 data block이  
멀리 떨어져



# First File System

- The file system would end up getting quite fragmented
  - free space was not carefully managed.
  - a logically contiguous file would be accessed by going back and forth across the disk, thus reducing performance dramatically.



- defragmentation tools can help
  - reorganize on-disk data to place files contiguously and make free space for one or a few contiguous regions, moving data around and then rewriting inodes and such to reflect the changes.

high overhead

- One other problem: the original block size was too small (512 bytes).
  - Smaller blocks minimized internal fragmentation (waste within the block),
  - but bad for transfer as each block/might require a positioning overhead to reach it.


작은 4KB 사용

Small block의 internal fragmentation은  
줄지만 data transfer에는 비효율적임

# FFS: Disk Awareness Is The Solution

→ hardware structure에 aware


- Fast File System (FFS)
  - Design the file system structures and allocation policies to be “disk aware” and thus improve performance
  - Same APIs, but advanced internal implementation
- FFS divides the disk into a number of cylinder groups.
  - A single cylinder is a set of tracks on different surfaces of a hard drive that are the same distance from the center of the drive
  - aggregates each N consecutive cylinders into group



cylinder group에 data를 head로  
읽어쓰기 위해 seek time을 위한 overhead

# FFS: Disk Awareness Is The Solution

- Modern drives *disk 내 physical인 공간이 어떻게 알 수 없음* do not export enough information for the file system to truly understand whether a particular cylinder is in use
- Modern file systems (such as Linux ext2, ext3, and ext4) instead organize the drive into block groups *때문에 cylinder group 대신 block group이라는 logical space를 사용*
  - A block group is just a consecutive portion of the disk's address space *physical하게 자른 것임으로 가려*
- By placing two files within the same group, FFS can ensure that accessing one after the other will not result in long seeks across the disk.




per-group inode bitmap (ib) and data bitmap (db)

# Linux File System

- **ext2/ext3/ext4**

- Block group

- Kernel tries to keep the data blocks belonging to a file in the same block group to reduce file fragmentation
    - All the block groups have the same size and are stored sequentially



# Policies: How To Allocate Files and Directories


- Keep related stuff together (keep unrelated stuff far apart)!!!
- FFS has to decide what is "related" and place it within the same block group

- Simple placement heuristics
  - Placement of a directory

- find the cylinder group with
- a low number of allocated directories (to balance directories across groups)
- a high number of free inodes (to subsequently be able to allocate a bunch of files)
- other heuristics could be used here (e.g., a high number of free data blocks)

## – File

- allocate the data blocks of a file in the same group as its inode
- places all files that are in the same directory in the cylinder group of the directory they are in.
- if a user creates four files, /a/b, /a/c, /a/d, and b/f, FFS would try to place the first three in the same group and the fourth in some other group.




same parent directory a  
(/a의 block group에)

다른  
group으로

# Policies: How To Allocate Files and Directories

## FFS policies

- (1) The data blocks of each file are near each file's inode
- (2) The files in the same directory are near one another



directory hierarchy

⇒ HDD → (seek)

same block group

SSD Random access 보인  
sequential access의 성능이 더 좋다.  
(locality이 기반한 allocation이 여전히 성능↑)

Are these policies still helpful  
for SSD-based systems?

Flash

→ Yes

| group | inodes | data    |
|-------|--------|---------|
| 0     | /      | /       |
| 1     | acde   | acddeee |
| 2     | bff    | bff     |
| 3     |        |         |
| 4     |        |         |
| 5     |        |         |
| 6     |        |         |
| 7     |        |         |
| ...   |        |         |

(1) & (2) (namespace-based locality)

| group | inodes | data |
|-------|--------|------|
| 0     | /      | /    |
| 1     | a      | a    |
| 2     | b      | b    |
| 3     | c      | cc   |
| 4     | d      | dd   |
| 5     | e      | ee   |
| 6     | f      | ff   |
| 7     |        |      |
| ...   |        |      |



이것은 disk access overhead를 줄여주는 함

Only (1)만 적용할 경우  
같은 directory의 file이 모여 있지 않음



# Measuring File Locality

- Let's see if indeed there is namespace locality
- Path difference
  - how "far away" file accesses were from one another in the directory tree
  - Same file accesses: 0
  - Different file in the same directory: 1
- SEER traces
  - about 7% of file accesses were to the file that was opened previously
  - nearly 40% of file accesses were to either the same file or to one in the same directory (distance  $\leq 1$ )
  - 25% of file accesses were to files that had a distance of two:
    - a set of related directories and consistently jumps between them (e.g., /proj/src/foo.c and proc/obj/foo.o)
    - FFS does not capture this type of locality



# The Large-File Exception

- Problem of FFS policies

- If a large file would entirely fill the block group it is first placed within, it prevents subsequent “related” files from being placed within this block group, and thus may hurt file-access locality.

| group | inodes                 | data                                                    |
|-------|------------------------|---------------------------------------------------------|
| 0     | /a- <del>bc</del> ---- | /aaaaaaaaa aaaaaaaaaa aaaaaaaaaa a- <del>bbbbbbbb</del> |
| 1     | -----                  | <del>bbb</del> ccc-cc -----                             |
| 2     | -----                  | -----                                                   |
| ...   |                        |                                                         |

locality를 해칠


- Large-File Exception

- The first twelve direct blocks are allocated into the first block group (48KB)
- The next “large” chunk of the file in another block group.
  - Each subsequent indirect block, and all the blocks it pointed to. (1024 blocks = 4 MB)
- If the chunk size is large enough, little seeking overhead

| group | inodes  | data        |
|-------|---------|-------------|
| 0     | /a----- | /aaaaa----- |
| 1     | -----   | aaaaa-----  |
| 2     | -----   | aaaaa-----  |
| 3     | -----   | aaaaa-----  |
| 4     | -----   | aaaaa-----  |
| 5     | -----   | aaaaa-----  |
| 6     | -----   | -----       |
| ...   |         |             |

# A Few Other Things About FFS

- Small files < 4KB
  - Internal fragmentation
  - sub-block allocation (512B)
  - if the file grows and requires a full 4KB of data, copy the sub-blocks into a 4KB block, and free the sub-blocks for future use.
  - To avoid a lot of extra I/O to perform the copy, the `libc` library would buffer writes and then issue them in 4KB chunks to the file system.
- Made the system more usable
  - Long file name (only 8 characters in the traditional fixed-size approach)
  - Symbolic link
  - Atomic `rename()`
- Many recent file systems take cues from FFS



EXT4 ↑  
2014-11-14

# Homework

---

- Homework in Chap 41 (FFS)

# Quiz

1. FFS divides the disk into cylinder groups to improve locality. (O)
2. In FFS, each cylinder group contains its own inode bitmap and data bitmap. (O)
3. In FFS, the data blocks of a file are placed ~~far from~~ the file's inode. (X)  
*같은 cylinder group → 데이터를 인접한 블록*
4. FFS tries to place files from the same directory into the same cylinder group. (O)
5. FFS places all blocks of a file in the same group ~~regardless of size~~. (X)  
*너무 크면 → 분산 저장*
6. FFS provides atomic rename operation. (O)
7. FFS tries to place related files in ~~different~~ block groups for parallel access. (X)
8. For small files (<4KB), FFS uses sub-block allocation to reduce internal fragmentation. (O)
9. Modern Linux file systems such as ext2/3/4 use ~~cylinder groups~~, just like FFS did. (X)  
*data group*
10. Since SSDs do not have seek latency, block groups are ~~not required~~ on SSDs. (X)  
*SSD에는 필요없음*